



Politecnico  
di Torino  
International  
University

**LAVAZZA**  
TORINO, ITALIA, 1895

# AI PERSONAS

Enrico Chen - s337750  
Van Thanh Nguyen - s336748  
Xiaoning Ma - s337332

# Project overview

## OUR GOAL

- Build an AI platform that allows internal business teams to interact with AI personas representing different market segments to explore and test strategies.

## KEY FEATURES

- **Automatically** organize unstructured data
- Data-driven AI personas based on **real** market data
- **Explainable** answers with clear references

## BENEFITS

- Reduce time and resources spent on data preprocessing
- Provide trustworthy and credible responses
- Enable more focused and informed market strategies

# VALUE PROPOSITION

For **business units** struggling in **evaluating marketing performances, customer understanding, models and ideas testing**, our platform allows **interacting with data-driven AI Personas** representing the different **market segments**



# Research Questions

## Extraction

How can we automatically extract accurate data from **complex, visually rich** PDFs while preserving layout semantics?

## Personality

Can **latent reasoning traits** (style profiles & value frames) be inferred from heterogeneous, noisy raw data to build consistent AI personas?

## Grounding

How can **hallucinations** in language model outputs be reduced?



# Research Question 1

## Question:

How can we automatically extract accurate data from **complex, visually rich** PDFs while preserving layout semantics?

## Our solution:

We treat PDFs as visual documents and extract semantic meaning using a VLLM



- Preserves layout and spatial relationships
- Fully automated, no manual rules
- Directly produces schema-consistent data ready for downstream use

## Alternative Solutions:

- **Manual annotation:** high accuracy, but low scalability, high costs
- **OCR-based pipelines:** scalable, but limited layout information

# Research Question 2

## Question:

Can we **infer latent reasoning traits** from heterogeneous, noisy raw data in order to build consistent AI personas?

## Our solution:

We use a cooperative multi-agent system in which each agent specialises in a specific stage of reasoning and trait inference.



Agent 1

Extract from raw data to get persona information

Agent 2

Infer latent persona reasoning traits by identifying patterns in language use, metrics, and expressed priorities

## Alternative Solutions:

- **Human review and inference:** high accuracy but slow, subjective, and not scalable
- **Single-model end-to-end inference:** automated, but sensitive to noise and lacking trait consistency

- ✓ Scalable and repeatable across datasets
- ✓ Reduces subjective human bias
- ✓ Produces consistent persona traits at scale

# Research Question 3

## Question:

How can we reduce **hallucinations** in language model outputs?

## Our solution:

We integrate a RAG system that retrieves relevant evidence before generation and constraints model outputs to the retrieved context.



- Improves factual accuracy
- Reduces hallucinations caused by unsupported generation
- Explainable responses
- Low computational costs

## Alternative Solutions:

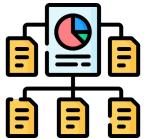
- **Prompt-based:** lightweight, but unreliable and sensitive to prompt design
- **Pure fine-tuning:** improves fluency, but does not guarantee factual grounding, computationally expensive

# Dataset



## Source & Scale

- Kantar France 2023
- 4,001 Respondents
- 9 consumer segments
- Rich qualitative + quantitative data



## Data Structure

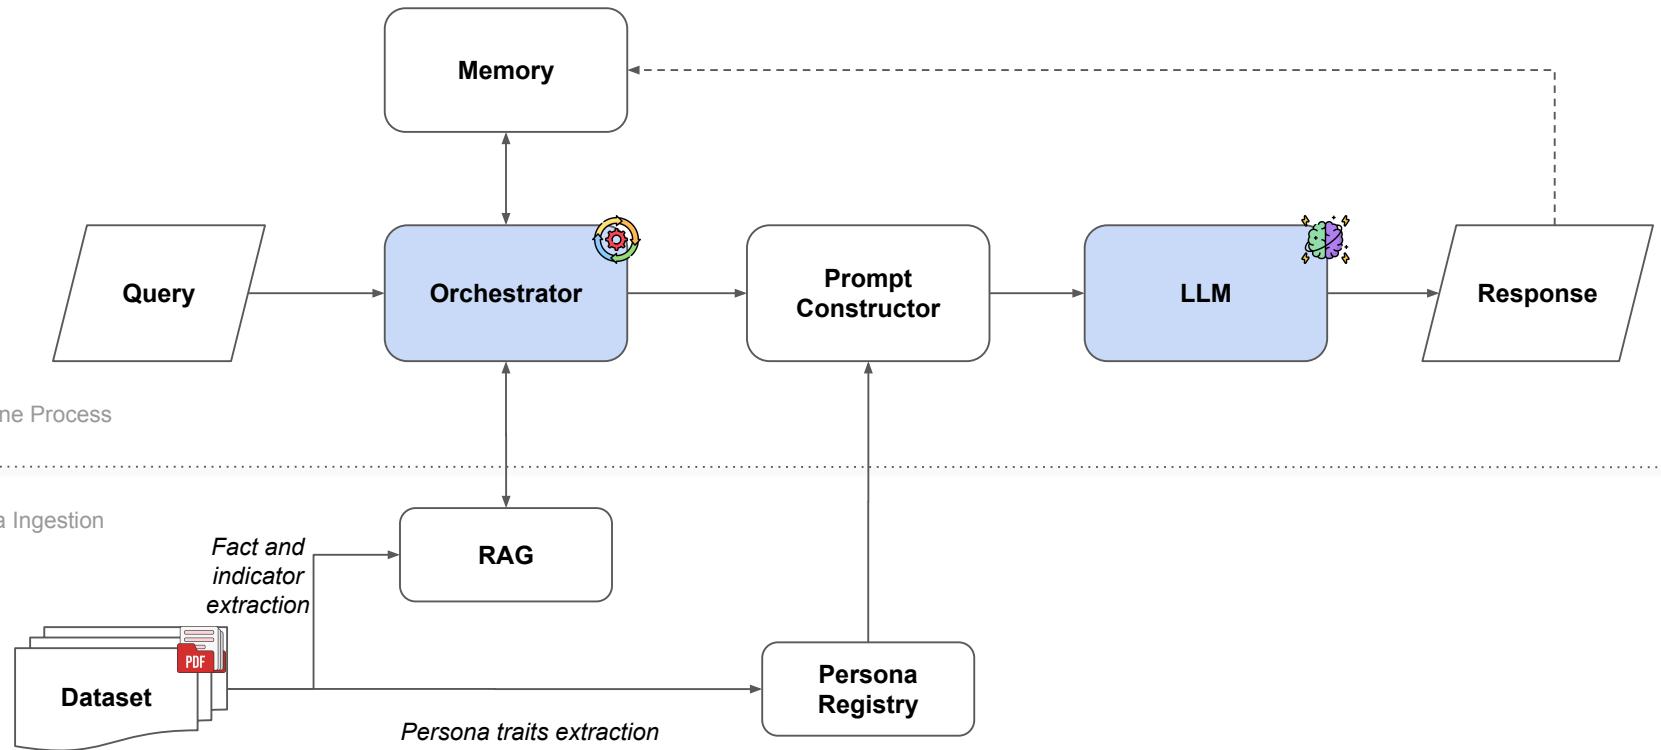
- Multi-modal composition
- Inconsistent layouts across segments
- Mixed granularity
- No unified or machine-readable format



## Key Dimensions

- Demographics
- Psychographics
- Coffee attitudes
- Consumption behaviours
- Brand perception & sustainability attitudes

# Functional Diagram



Fact data: all data (including segments characteristics, market data)

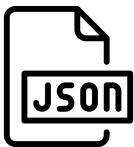
Persona traits data: characteristics of the customer segments (like tone, style, values, preferences)

# Method - Persona Extraction



Each page of the dataset

VLLM Extraction



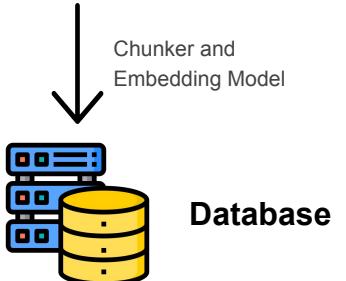
VLLM Reasoning



Prompt tuning



JSON indicators



Persona Traits

Database

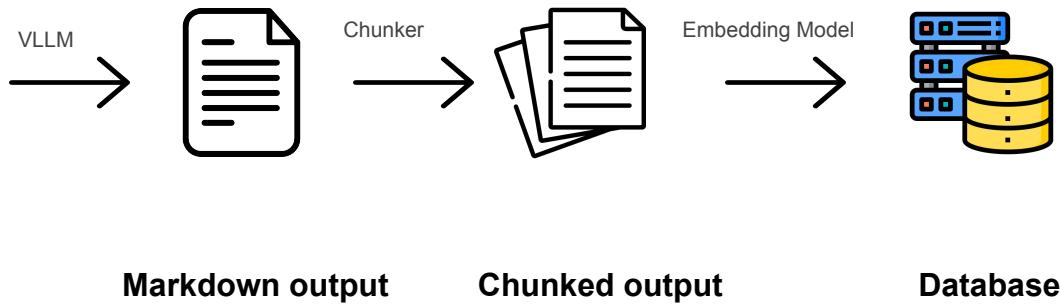
Chunker and  
Embedding Model

Fine-tuned AI  
Personas

# Method - Fact Extraction



Each page of the dataset

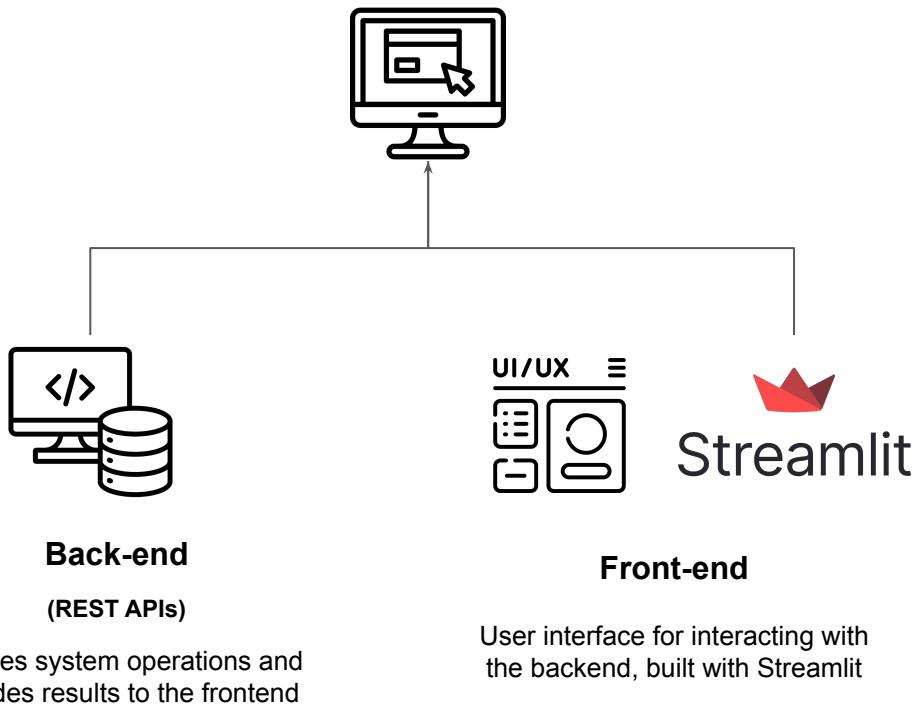


Markdown output

Chunked output

Database

# Method - Application



## Features implemented

- 1 Authentication
- 2 Customer segment selection
- 3 Past chat selection
- 4 Context customization
- 5 Persona name customization
- 6 Persona information

# UI - Login

Deploy :

## Lavazza AI Personas

---

### Authentication

> API Configuration

API server is online

Login Register

---

Username

Access Token

# UI - Registration

Deploy :

## Lavazza AI Personas

---

### Authentication

> API Configuration

API server is online

Login [Register](#)

Username

Access Token  
 

Confirm Token  
 

[Register](#)

# UI - Homepage

Welcome, Alice

[Logout](#)

---

Select Persona

- [o Conscious Explorers](#)
- [➤ About Conscious Explorers](#)
- [o Curious Connoisseurs](#)
- [➤ About Curious Connoisseurs](#)
- [o Habitual Disciplinists](#)
- [➤ About Habitual Disciplinists](#)
- [o Juvenile Dabblers](#)
- [➤ About Juvenile Dabblers](#)
- [o Matured Loyalists](#)
- [➤ About Matured Loyalists](#)
- [o Mindful Engagers](#)

Please select a persona from the sidebar to start chatting

Lavazza AI Personas

Deploy :

# UI - Persona name customization

Welcome, Alice

[Logout](#)

---

Select Persona

- Conscious Explorers
- About Conscious Explorers
- Curious Connoisseurs
- About Curious Connoisseurs
- Habitual Disciplinists
- About Habitual Disciplinists
- Juvenile Dabblers
- About Juvenile Dabblers
- Matured Loyalists
- About Matured Loyalists
- Mindful Engagers

---

Lavazza AI Personas

Start Chat with Curious Connoisseurs

Give your assistant a custom name for this session:

Assistant Name

 ⓘ

[Start Chat](#) [Cancel](#)

# UI - Chat Management

- o Matured Loyalists
- > About Matured Loyalists
- o Mindful Engagers
- > About Mindful Engagers
- o Routine Rechargers
- > About Routine Rechargers
- o Spontaneous Connectors
- > About Spontaneous Connectors
- o Uninvolved Pragmatists
- > About Uninvolved Pragmatists

---

**Chat Sessions**

- ✓ Giulia (0) Delete

---

- > API Info

**Lavazza AI Personas**

💬 Chat
👤 Persona Info
Deploy
⋮

---

Chatting with: Giulia Clear Chat

---

Start the conversation by typing a message below!

---

▼ Chat Settings

Show retrieved context  ⓘ
Number of context documents (top-k)

 
15

---

Type your message:

Ask me anything about coffee...

Send

New Chat

# UI - Chat Interaction

Deploy :

- o Matured Loyalists
- > About Matured Loyalists
- o Mindful Engagers
- > About Mindful Engagers
- o Routine Rechargers
- > About Routine Rechargers
- o Spontaneous Connectors
- > About Spontaneous Connectors
- o Uninvolved Pragmatists
- > About Uninvolved Pragmatists

**Chat Sessions**

- ✓ Giulia (2) 

> API Info

## Lavazza AI Personas

Chat Persona Info

Chatting with: Giulia Clear Chat

You:  
Hi who are you?

Assistant:  
I'm Giulia, a coffee specialist dedicated to exploring high-quality, innovative, and sustainable coffee experiences.

Retrieved Context

Chat Settings

Show retrieved context  Number of context documents (top-k) 15

# Evaluation

## Persona Extraction

Is persona information correctly extracted?



## Fact Extraction

Is extracted fact data correct?

## Retrieval Relevance

Is the retrieved information relevant?

## Authenticity Evaluation

Do personas reflect real consumer segments?

# Evaluation - Persona Extraction

## *Approach*



Compare persona extraction outputs with manually annotated ground truth.

## *Metrics*

- **Persona Detection Rate**
- **Metrics Recall**
- **Metrics Precision**

Are all target personas identified?  
Is the required information complete?  
Is the extracted information accurate?

\*Persona Detection Rate = (Correctly Detected Personas) / (Total Ground Truth Personas)

\*Metrics Recall = (Correctly Metrics) / (Total Ground Truth Metrics)

\*Metrics Precision = (Correct Metrics) / (All Extracted Metrics)

# Evaluation - Persona Extraction

## Configuration

- **Models:**
  - ◆ **Extract indicators:** mistralai/mistral-medium-3-instruct
  - ◆ **Persona traits reasoning:** mistralai/mistral-medium-3-instruct
- **Test dataset:**
  - ◆ **Scope:** 23 pages focused on the Curious Connoisseurs segment
  - ◆ **Source:** Customer Segmentation Analysis PDF
  - ◆ **Ground Truth:** 1,051 metrics manually extracted and validated from the PDF

## Result

	Score	Interpretation
Persona Detection Rate	100%	All personas correctly identified (23/23)
Metrics Recall	95.30%	Very high coverage of ground-truth metrics (1002/1051)
Metrics Precision	96.80%	Minimal noise in extracted metrics (1002/1035)

# Evaluation - Fact Extraction

## Approach



- Manually verify numeric and factual data against source PDFs.
- Run extraction and check exact matches.
- A result is correct only if  $\geq 80\%$  textual overlap with the extracted markdown.

## Metrics

→ **Exact Match Accuracy**

what fraction of expected facts were extracted correctly?

\*Exact Match Accuracy = (Exactly Matching Values) / (Total Ground Truth Values)

# Evaluation - Fact Extraction

## Configuration

- **Models:**
  - ◆ **PDF to Markdown Conversion:** mistralai/mistral-large-3-675b-instruct-2512
- **Test dataset:**
  - ◆ **Scope:** 23 pages focused on the Curious Connoisseurs segment
  - ◆ **Source:** Customer Segmentation Analysis PDF
  - ◆ **Ground Truth:** 467 validated metric and statement snippets manually extracted from the PDF

## Result

### Score

Exact Match Accuracy

97%

### Interpretation

High extraction accuracy with minimal deviation from ground truth

# Evaluation - Retrieval Relevance



## Approach

- Create persona-specific test queries.
- Retrieve top-K documents ( $K = 3, 5, 10, 20$ ..)
- Manually label relevance for retrieval documents.

## Metrics

- **Precision@K**      How many of the top K results are actually correct.
- **Recall@K**      How many of the total correct results were successfully found within the top K.

\*Precision@K = (Relevant docs in top-K) / K

\*Recall@K = (Relevant docs in top-K) / (Total relevant docs)

# Evaluation - Retrieval Relevance

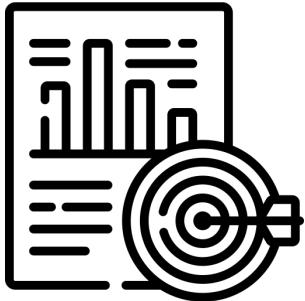
## Configuration

- **Models:**
  - ◆ **Embedding:**  
sentence-transformers/all-mnlp-base-v2
- **RAG setup:**
  - ◆ **Chunk Size:** 1,200 characters
  - ◆ **Chunk Overlap:** 50 characters
  - ◆ **Input:** 222-page split to 712 total text chunks
- **Test dataset:**
  - ◆ **Source:** Customer Segmentation Analysis PDF
  - ◆ **Ground Truth:** 31 evaluation questions focused on the Curious Connoisseurs segment

	Score	Interpretation
Precision@3	58.10%	Top-3 results are moderately relevant
Precision@5	56.80%	Relevant remains stable in top-5
Precision@10	54.20%	Slight precision drop
Precision@20	55.50%	Precision stabilizes at broader context window
Recall@3	16.70%	Limited coverage with very short context
Recall@5	24.30%	Partial retrieval of relevant context
Recall@10	45.10%	Balanced trade-off between precision and recall
Recall@20	94.10%	Near-complete retrieval of relevant

**Result**

# Evaluation - Persona Authenticity



## Approach

- Generate persona-based responses to predefined questions.
- Have domain experts score them on authenticity, style alignment, and factual grounding, each on a 1–5 scale.

## Metrics

- |                                    |  |
|------------------------------------|--|
| → <b>Expert Authenticity Score</b> | Average authenticity ratings.          |
| → <b>Style Alignment Score</b>     | Average all style alignment ratings.   |
| → <b>Factual Grounding Score</b>   | Average all factual grounding ratings. |

# Evaluation - Authenticity Results

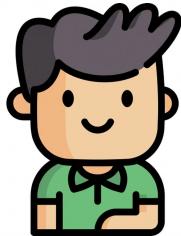
## Configuration

- **Models:**
  - ◆ **Filter context model:** mistralai/mistral-medium-3-instruct
  - ◆ **Generation model:** mistralai/mistral-small-24b-instruct
- **RAG Setup:**
  - ◆ **Chunk Size:** 1,200 characters (~300–350 tokens)
  - ◆ **Chunk Overlap:** 50 characters (to preserve context continuity)
  - ◆ **Input:** 222-page to 712 total text chunks
- **Test dataset:**
  - ◆ **Source:** Customer Segmentation Analysis PDF
  - ◆ **Ground Truth:** 31 evaluation questions focused on the Curious Connoisseurs segment

## Result

	Score	Interpretation
<b>Expert Authenticity Score</b>	3.90/5	Persona behavior is largely authentic.
<b>Style Alignment Score</b>	3.74/5	Style is mostly consistent, with minor persona drift.
<b>Factual Grounding Score</b>	3.67/5	Responses are generally grounded.

# Manage



## Enrico

- Semantic extraction of fact data
- Fact data indexing and organization
- Frontend application development
- Related presentation slides



## Thanh

- Persona semantic information extraction
- Backend system development
- Evaluation (persona & fact data extraction)
- Related presentation slides

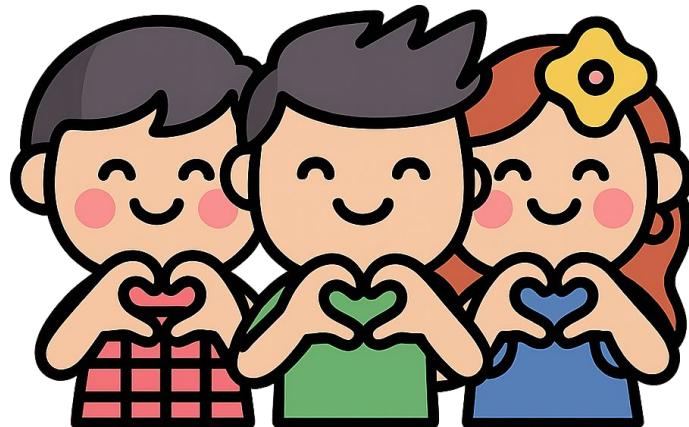


## Xiaoning

- AI persona common trait extraction
- Evaluation dataset construction
- Evaluation (retrieval & authenticity analysis)
- Related presentation slides



**Politecnico  
di Torino**  
International  
University



**LAVAZZA**  
TORINO, ITALIA, 1895

 **THANK YOU**

**Enrico Chen - s337750**  
**Van Thanh Nguyen - s336748**  
**Xiaoning Ma - s337332**



# Appendix

# Manage - Gantt

WB No	TASK TITLE	DEPEND ON	OWNER	COLLABORATORS	START	END	PERSON WEEK	PROGRESS (%)	PHASE			
									11/2025			
					w1	w2	w3	w4	w5	w6	w7	w8
1	Project Conception and Initiation		Thanh	Others	03/11/2025	14/11/2025	3.00					
1.1	Project structure+ work flow		Thanh	Others	03/11/2025	03/11/2025	0.50	100.00				
1.2	Kick off with Lavazza tutor		XiaoNing	Others	03/11/2025	07/11/2025	0.50	100.00				
1.3	Research		Enrico	Others	03/11/2025	14/11/2025	2.00	100.00				
2	Design	1	Enrico	Others	03/11/2025	16/11/2025	7.50					
2.1	Objective and Goal Definition	1.2	Thanh	Others	03/11/2025	14/11/2025	0.50	100.00				
2.2	Stakeholder Map	1.2	Enrico	Others	03/11/2025	14/11/2025	1.00	100.00				
2.3	User Personas Definition	1.2	Enrico	Others	03/11/2025	14/11/2025	0.50	100.00				
2.4	User Journey Definition	1.2	XiaoNing	Others	03/11/2025	14/11/2025	1.00	100.00				
2.5	User Requirements Definition	1.2	XiaoNing	Others	03/11/2025	14/11/2025	0.50	100.00				
2.6	Usecase diagram	1.2;1.3	XiaoNing	Others	03/11/2025	16/11/2025	1.00	100.00				
2.7	Func and Non-Func Requirements Definition	1.2;1.3	Enrico	Others	03/11/2025	16/11/2025	1.00	100.00				
2.8	System Architecture and Func Diagram	1.2;1.3	Thanh	Others	03/11/2025	16/11/2025	1.00	100.00				
2.9	Risk Analysis	1.2;1.3	Thanh	Others	08/11/2025	16/11/2025	1.00	100.00				
3	Management	2	XiaoNing	Others	08/11/2025	16/11/2025	1.00					
3.1	Tasks Breakdown and Gantt Diagram	2	XiaoNing	Others	08/11/2025	16/11/2025	1.00	100.00				



# Manage - Gantt

WB No	TASK TITLE	DEPEN ON	OWNER	COLLABORATORS	START	END	PERSON WEEK	PROGRESS (%)	PHASE ONE															
									11/2025				12/2025											
									w1	w2	w3	w4	w5	w6	w7	w8	w9	w10	w11	w12	w13	w14	w15	w16
4	Data Foundation	1.2	Thanh	Others	17/11/2025	23/11/2025	6.00																	
4.1	Data acquisition & ingest	1.2	Thanh	Others	17/11/2025	23/11/2025	3.00	90.00																
4.1.1	Collect data from Lavazza	1.2	Enrico	Others	17/11/2025	21/11/2025	1.00	100.00																
4.1.2	Collect data from external source		Thanh	Others	17/11/2025	21/11/2025	1.00	100.00																
4.1.3	Understand dataset	4.1.1;4.1.2	XiaoNing	Others	21/11/2025	23/11/2025	1.00	100.00																
4.2	Finalize PersonaProfile schema	4.1.3	XiaoNing	Others	20/11/2025	28/11/2025	1.00	100.00																
4.3	Data processing pipelines	4.2	Thanh	Others	20/11/2025	30/11/2026	2.00	100.00																
4.3.1	Handle fact data pipeline	4.3	Enrico	Thanh	20/11/2025	30/11/2025	1.00	100.00																
4.3.2	Handle persona data pipeline	4.3.1	XiaoNing	Thanh	20/11/2025	30/11/2025	1.00	100.00																
5	Prompt Tuning AI Persona	4.3.2	XiaoNing	Thanh	24/11/2025	27/12/2025	4.00																	
5.1	Semantic Extraction/Personas Structuring from Customer Segmentation Data	5	XiaoNing	Thanh	24/11/2025	12/12/2025	2.00	100.00																
5.1.1	Extract common traits/rules for each personas	5.1	XiaoNing	Thanh	24/11/2025	12/12/2025	1.00	100.00																
5.1.2	Create personas fine-tuning dataset	5.1.1	XiaoNing	Thanh	30/11/2025	12/12/2025	1.00	100.00																
5.2	Implement Training pipeline	5.1.2	Thanh	XiaoNing	01/12/2025	28/12/2025	1.50	30.00																
5.3	Implement Inference & serving	5.2	Thanh	XiaoNing	13/12/2025	28/12/2025	0.50	0.00																
6	Fact Data Ingestion	5.3	Enrico	Others	24/11/2025	21/12/2025	2.50																	
6.1	Semantic Extraction from Fact Data	6	Enrico	Thanh	24/11/2025	21/12/2025	1.00	50.00																
6.2	Indexing fact data	6.1	Enrico	Thanh	24/11/2025	21/12/2025	0.50	30.00																
6.3	Implement retrieval logic with RAG	6.2	Enrico	Thanh	01/12/2025	21/12/2025	1.00	30.00																



# Manage - Gantt

WB No	TASK TITLE	DEPEND ON	OWNER	COLLABORATORS	START	END	PERSON WEEK	PROGRESS (%)	PHASE ONE										
									11/2025				12/2025				01/2026		
									w1	w2	w3	w4	w5	w6	w7	w8	w9	w10	
7	Core Layer	4;5;6	Thanh	Others	01/12/2025	28/12/2025	8.50												
7.1	Input handling	4	XiaoNing		01/12/2025	14/12/2025	1.00	20.00											
7.2	Retrieval-Augmented Generation	6	Enrico		07/12/2025	28/12/2025	1.50	20.00											
7.3	Implement Orchestrator logic	7.1;7.2	Thanh	Others	07/12/2025	28/12/2025	2.00	10.00											
7.4	Implement Explanation module	7.3	Thanh		07/12/2026	28/12/2025	1.00	0.00											
7.5	Prompt construction	7.3	Enrico	Others	07/12/2027	28/12/2025	1.00	10.00											
7.6	Persona registry	5	XiaoNing		07/12/2025	28/12/2025	1.00	20.00											
7.7	AI Persona Router	5	Thanh		07/12/2025	28/12/2025	1.00	0.00											
8	Application Layer	7	XiaoNing	Others	15/12/2025	04/01/2026	2.00												
8.1	Persona configuration	7.6	Thanh	Others	15/12/2025	04/01/2026	1.00	0.00											
8.2	Q&A service	7	Enrico	Others	15/12/2025	04/01/2026	1.00	0.00											
9	UI	8	XiaoNing	Others	22/12/2025	04/01/2026	1.50												
9.1	FE	8	XiaoNing	Others	22/12/2025	04/01/2026	1.50	0.00											



# Manage - Gantt

WB No	TASK TITLE	DEPEN ON	OWNER	COLLABORATORS	START	END	PERSON WEEK	PROGRESS (%)	PHASE ONE									
									11/2025		12/2025		01/2026		02/2026		03/2026	
w1	w2	w3	w4	w5	w6	w7	w8	w9	w10	w11	w12	w13	w14	w15	w16	w17	w18	
10	Monitoring and Evaluation	7	Enrico	Others	22/12/2025	11/01/2026	2.00	0.00										
10.1	Evaluation	7	Enrico	Others	22/12/2025	11/01/2026	2.00	0.00										
11	Deployment	7;8;9	Thanh	Others	22/12/2025	11/01/2026	2.00	0.00										
11.1	Packaging	7;8;9	Thanh	Others	22/12/2025	11/01/2026	1.00	0.00										
11.2	Deploy entire system	11.1	Thanh	Others	22/12/2025	11/01/2026	1.00	0.00										
12	Testing	7;8;9	XiaoNing	Others	22/12/2025	11/01/2026	2.00	0.00										
12.1	Test	7;8;9	XiaoNing	Others	22/12/2025	11/01/2026	1.00	0.00										
12.2	Fix Bug	12.2	Thanh	Others	22/12/2025	11/01/2026	1.00	0.00										
13	Demo	12	Enrico	Others	29/12/2025	11/01/2026	1.00	0.00										
13.1	Run full flow & get feedback	12	Enrico	Others	22/12/2025	11/01/2026	1.00	0.00										



# Manage - Gantt

WB No	TASK TITLE	DEPENDS ON	OWNER	COLLABORATORS	START	END	PERSON WEEK	PROGRESS (%)	PHASE ONE										
									11/2025		12/2025		01/2026		02/2026		03/2026		04/2026
w1	w2	w3	w4	w5	w6	w7	w8	w9	w10	w11	w12	w13	w14	w15	w16	w17	w18	w19	
14	Communication		Enrico	Others	14/11/2025	19/11/2025	5.00												
14.1	First Checkpoint Presentation		Enrico	Others	14/11/2025	19/11/2025	1.00	100.00											
14.2	Second Checkpoint Presentation		XiaoNing	Others	02/12/2025	09/12/2025	1.00	100.00											
14.3	Third Checkpoint Presentation		Thanh	Others	30/12/2025	06/01/2026	1.00	0.00											
14.4	Final Presentation		Enrico	Others	23/12/2025	12/01/2026	1.00	0.00											
14.5	Final Report		XiaoNing	Others	23/12/2025	13/01/2026	1.00	0.00											



# Design

## 1. User Interface (UI)

The user interface serves as the system's entry point, built as a **Frontend (FE)** application. It enables users to interact seamlessly with the platform, submit queries, upload data, and view results or reports.

## 2. Application Layer

This layer contains the core application logic and manages all user-driven workflows.

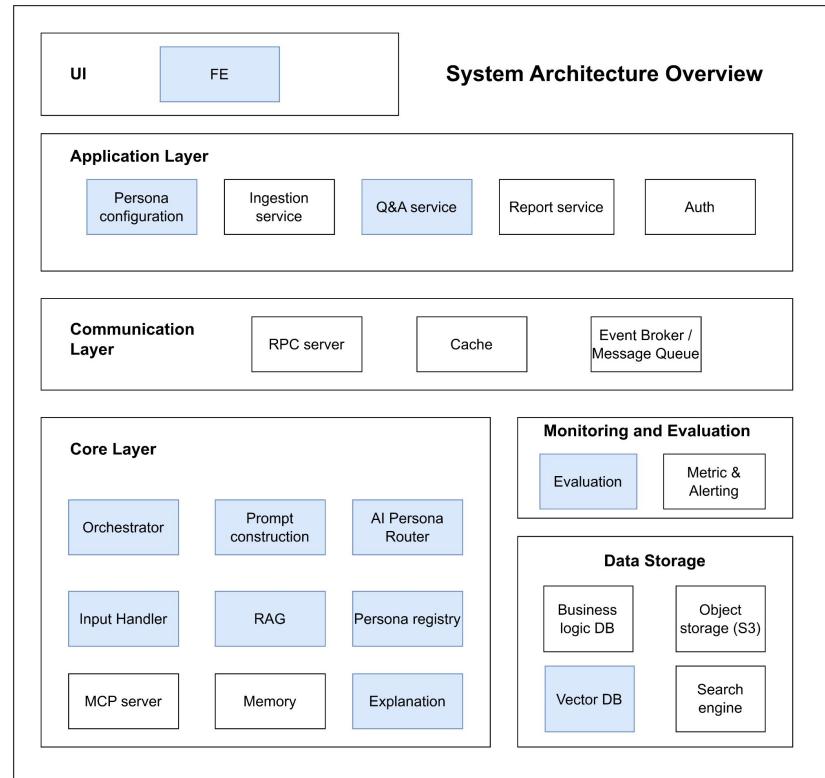
Key components include:

- **Persona Configuration:** Enables users to select or customize AI personas dynamically.
- **Ingestion Service:** Handles ingestion of raw data such as PDFs or images and stores them in S3.
- **Report Service:** Generates structured, formatted reports from processed and analyzed data.
- **Q&A Service:** Manages interactive question-and-answer exchanges with the AI.
- **Auth Service:** Provides authentication and authorization for users, ensuring secure access and operations.

## 3. Communication Layer

This layer facilitates efficient communication and coordination among microservices.

- **RPC Server:** Enables direct service-to-service communication via Remote Procedure Calls.
- **Cache:** A high-speed memory layer that stores frequently accessed data to optimize performance.
- **Event Broker / Message Queue** (RabbitMQ or Kafka): Handles asynchronous communication and event-driven processing across services, ensuring reliability, scalability, and robust monitoring.

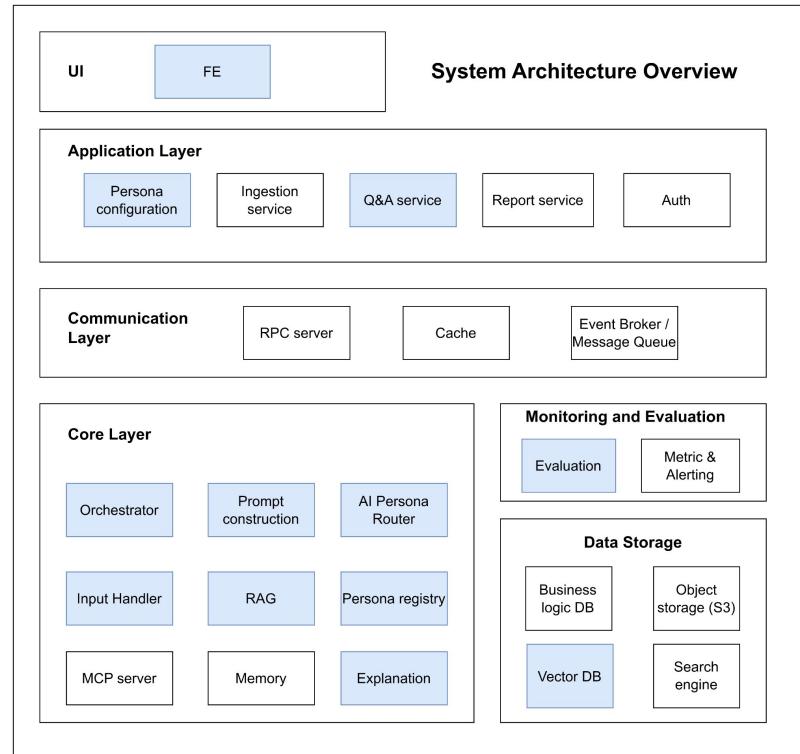


# Design

## 4. Core Layer

The intelligence engine of the system—handles AI persona logic, LLM orchestration, and data-driven grounding.

- **Orchestrator:** The central coordinator of the Core Layer. When a request arrives, the Orchestrator manages the entire generation process, directing which services to call.
- **Input Handler:** Preprocesses and normalizes user inputs, including text extraction from PDFs and preparation of image data for AI analysis.
- **Prompt Construction:** Dynamically builds structured prompts by combining user input, persona rules, and retrieved data.
- **AI Personas:** Represents the fine-tuned Large Language Models (LLMs) tailored to embody distinct customer segment personalities.
- **RAG (Retrieval-Augmented Generation):** Provides factual grounding by retrieving relevant information from the Vector DB, ensuring responses remain accurate.
- **Persona Registry:** Stores the static attributes and behavioral definitions of each persona, guiding prompt construction and response tone.
- **Explanation:** This module allows for an in-depth explanation of the thought process behind the reasoning model and the data used in the thinking process.
- **MCP Server (Model Context Protocol Server):** Enriches LLM interactions with real-time contextual or external domain data.
- **Memory:** It stores the recent history of the user's chat, allowing the persona to remember what was said earlier in the conversation and provide context-aware answers.



# Design

## 5. Monitoring and Evaluation

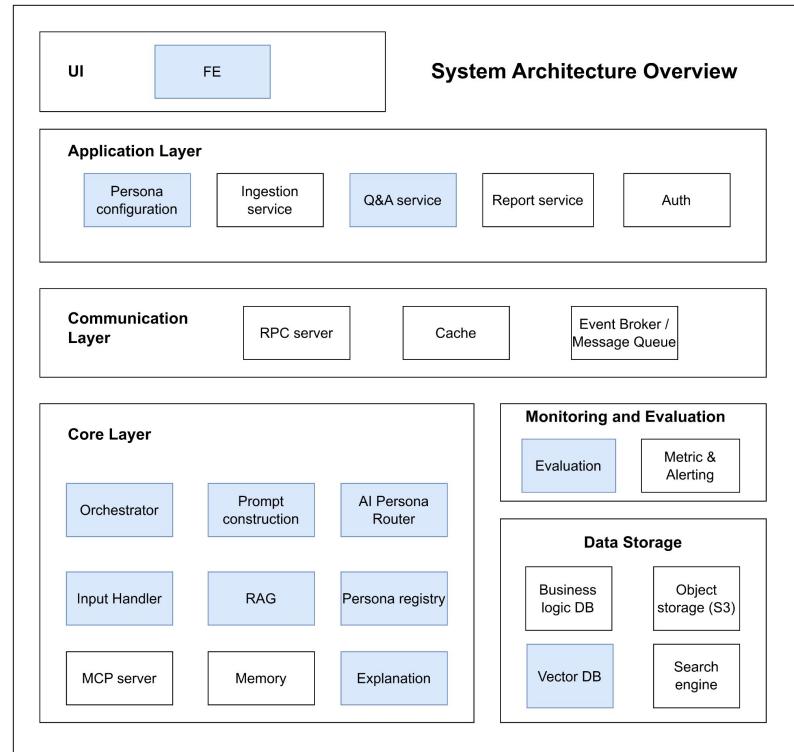
A centralized observability layer that tracks performance, quality, and reliability across all services.

- **Evaluation Tools:** Measure the accuracy and quality of AI responses and data processing outcomes.
- **Metrics & Alerting:** Monitor key indicators such as latency, error rates, resource utilization, and token usage, triggering alerts for anomalies or system degradation.

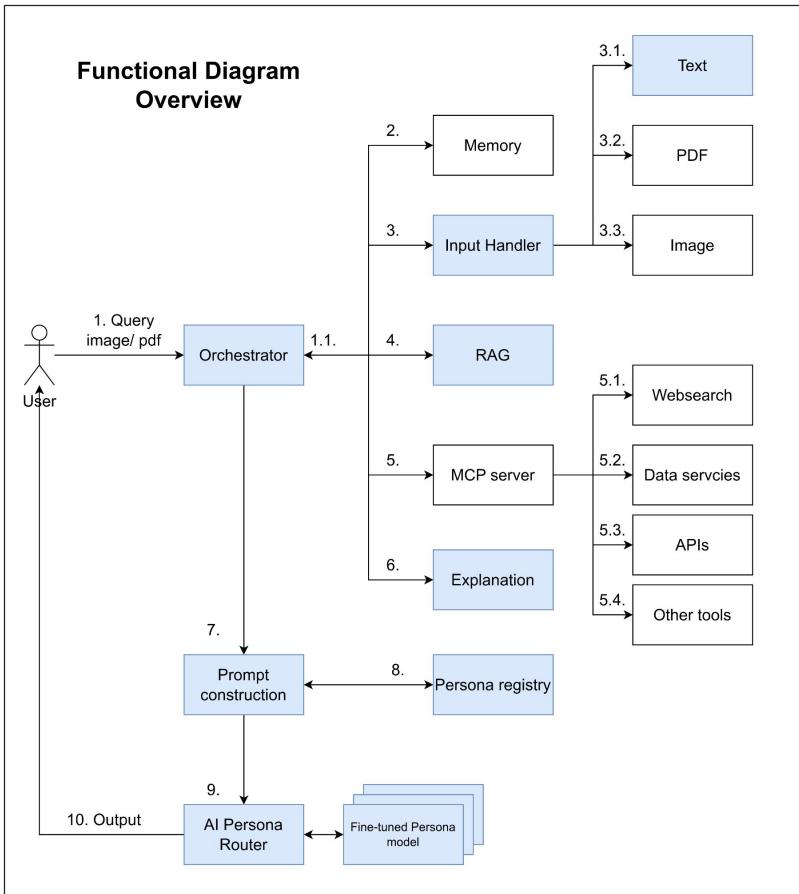
## 6. Data Storage Layer

The persistence foundation of the system, designed for scalability, durability, and speed.

- **Business Logic Database:** Stores structured data such as user profiles, authentication records, saved reports, and persona definitions.
- **Object Storage (S3):** Manages large, unstructured data files (e.g., raw PDFs, images, and uploaded datasets).
- **Vector Database:** Stores embeddings for persona-related documents, historical interactions, and reference materials — powering RAG retrieval and factual grounding.



# Design



**1. User Query Submission:** User sends a query with optional attached files (image, PDF, etc.) to the Orchestrator.

**1.1 Orchestrator Analysis:** The Orchestrator analyzes the query and attachments to decide which services should be used.

**2. Memory Integration :** Extract useful information from chat history.

**3. Input Preprocessing:** Inputs are preprocessed before passing to the model.

**3.1 Text Input:** Normalize text to make it easier to handle in later steps.

**3.2 PDF Input:** Parse, process, and extract meaningful information from PDF files.

**3.3 Image Input:** Process images and extract valuable information.

**4. Context Retrieval (RAG System):** Use the query and relevant input information to retrieve context (e.g., market data) via a RAG system.

**5. Tool Selection & MCP Server Requests**

- Decide which tools should be used to enrich the context.
- Send requests to the MCP server to gather corresponding context.

**5.1 Web Search:** Extract updated information from the internet (trends, real-time data, missing internal data, etc.).

**5.2 Database Query:** Retrieve useful data from internal or external databases.

**5.3 External APIs:** Call APIs to obtain additional information.

**5.4 Other Tools:** Use calculators, simulators, weather data extractors, or other utilities to enrich context.

**6. Explanation:** The explanation module will explain in detail the thought process of the reasoning model and the data used for the thinking process.

**7. Prompt Construction:** The Orchestrator aggregates useful context and passes it to Prompt Construction.

**8. Persona Selection :** Apply the selected Persona profile, including: Demographics, Behavior Data, Transactional Data, ...

**9. Persona Model Routing**

- Route to a fine-tuned Persona model.
- Pass the enriched prompt and context.

**10. Model Response:** Generate a response with: Specific personality, Tone, Linguistic style of the Persona

# AI Personas Extraction & Fine-Tuning

## Indicators (VLLM extraction output)



### Example indicator shape (JSON):

```
{
  "Indicator": {
    "sources": {
      "url": "https://www.pclavazza.com/nead/20/Indicator"
    },
    "statements": {
      "statement": "individual insights within an indicator, sostons and influences",
      "metrics": "marks whethr a svisually emphasized: 'index', '%', 'count', 'rank'",
      "influence": "flags whether a statement shapes tone or stance, 'smex': 'sources'"
    }
  }
}
```



# AI Personas Extraction & Fine-Tuning

## Traits (reasoning output)

### Persona Blueprint (Traits)



#### style\_profile

⌚ how they speak: tone, formality, directness, emotional flavour, criticality, verbosity, preferred structures, example phrases.



#### value\_frame

⚖️ what they prioritize: priority\_rank (sustainability, price, etc.), novelty seeking, brand loyalty, health concern, description.



#### reasoning\_policies

⌚ purchase\_advice, product\_evaluation, information\_processing, content\_filters (biases, rules, praise/criticism triggers, trust, disclaimers).

### Example JSON Structure

• • •

```
// How the model should "speak"
style_profile: {
  tone_adjectives: string[], // ["Curious", "confident", "quality-focused", "pragmatic", ...],
  formality_level: "low" | "medium" | "high",
  directness: "very_direct" | "balanced" | "hedged",
  emotional_flavour: "neutral" | "enthusiastic" | "cool_detached" | "warm_reflective",
  criticality_level: "high" | "medium" | "low",
  verbosity_preference: "concise" | "detailed" | "varies_by_question",
  preferred_structures: string[], // ["bullet_point", "clear_rationale", "pros_cons", "step_by_step"]
  typical_register_examples: string[] // short example phrases in target style
},
```

```
// What they care about - used to bias recommendations / reasoning
value_frame: {
```

```
  priority_rank: string[] // e.g. ["quality", "convenience", "sustainability", "price"],
  sustainability_orientation: "high" | "medium" | "low",
  price_sensitivity: "high" | "medium" | "low",
  price_sensitivity: "high" | "medium" | "low",
  novelty_seeking: "high" | "medium" | "low"
},
```



# Design - Risks Analysis

## Technical Risks

- Hallucinations and inaccurate responses: mitigate with RAG system
- Insufficient critical thinking: mitigate with RAG and prompt engineering
- Opacity: mitigate with RAG
- Inconsistent or generic personality: mitigate by fine-tuning (in case of limited resource use PEFT, smaller models, RAG with few-shot prompting)
- Performance evaluation difficulty



# Design - Risks Analysis

## Governance and Security Risks

- Privacy and compliance with AI Act and GDPR
- Proprietary data protection
- System integration difficulty with existing systems and infrastructure



# Design - Risks Analysis

## Data and Other Risks

- Data integration difficulty
- Data quality and bias
- Over relying on AI Personas

